

# Assignment 3: HyperLogLog

Gustav Gyrst and Katarzyna Toborek

2021-11-03

## 1 Introduction

The following report provides a description of an implementation of **HyperLogLog**, which is an efficient algorithm used for estimating a number of distinct elements in a given input stream. The implementation makes use of a few smaller functions, implemented and tested independently before assembling the algorithm. The functions, described in more detail in the next section, are **hash**, **rho** and **registers**. Based on the registers, **HyperLogLog** produces an estimate of the stream cardinality. The total space usage to produce the result is  $m \log(\log n)$  bits, where  $m$  is the number of registers and  $n$  is the size of the input stream.

## 2 Implementation

All code parts for the project was implemented in Python 3.8.8.

### 2.1 Hash

The **hash** algorithm used for **hyperLogLog** is implemented using bitwise operations. A helper function is used that returns the number of set bits in a given input value. We use the given matrix  $A$  containing  $32 \times 32$  – *bit* integers to calculate the hash. For each value  $n$  from the input stream, the helper function is applied to the result of multiplying  $n$  with a single element from  $A$  (bitwise  $A[i] \& n$ ). Then the operation  $\&1$  is applied to the result, what corresponds to modulo 2 operation - if it evaluates to 0, the integer is even, if it's 1, the integer is odd. Next the result is multiplied by  $i$ th power of 2, what is done by applying a left shift  $i$  to the result ( $<< i$ ). These operations are done in a **for** loop, with  $i$  being in range  $0 - 32$ . With each iteration the results are accumulated to return a final hash value of  $n$ .

## 2.2 Rho

The `Rho` method returns the position of the first 1 in the binary representation. The bitwise right shift was used where we start with shifting  $32 - i$  (where  $i$  iterates over range 33) and then checks whether the current bit in this position is a 1 by using the bitwise AND operator. It returns  $i$  when it finds the first bit, otherwise it returns `None`.

## 2.3 Registers

The `registers` function combines both the `hash` function, `rho` and also introduces another hash-function  $f(x)$ , which purpose is to calculate an index for each input value. This index is used to map a value to certain "bin" in the register  $M - M[index]$ . When  $m$  is of size 1024, we use a right shift of 21 to get indexes between 0 – 1023. When we have the index for our input, we calculate the hash value of it using our `hash` function. Then we apply the `rho` function on this hash value and store the rho value for the hashed input under in the index given by  $f(x)$  in our register. If the current value at this position is smaller than the new value found, the current value is then replaced. The function then returns  $M$  with its values.

## 2.4 HyperLogLog

At last, `hyperLogLog` algorithm uses the results provided by `registers` to produce the desired output. The implementation follows the pseudocode provided in the problem description. We use the provided "magic" constant  $\alpha_m = 0.7213/(1 + 1.079/m)$  for the raw estimate and calculate the number of empty registers. Based on these values the final estimate is returned: for smaller cardinalities and at least one non-empty register linear counting is used; for larger estimates, large range correction is applied instead.

## 3 Testing for correctness

For the greater part our implementation was tested for correctness on Codejudge, where each function used for `HyperLogLog` was first tested independently before testing the algorithm performance. We assume that the quality of the Codejudge tests was much higher than anything we could come up with ourselves.

### 3.1 CodeJudge Tests

Each part of the implementation (`hash`, `rho`, `registers`, and `HyperLogLog`) was tested on different small and large inputs including some corner cases. For instance the `hash` algorithm was tested on small input the smallest and largest number (00000000 and ffffffff) and also a large test with an input of size 500+. Similarly for `rho` and `registers` the tests included both small and large inputs. In particular the `registers` implementation was tested on two large inputs.

`HyperLogLog` algorithm was tested with the *Threshold* exercise on Codejudge. There the success of the tests depended on the algorithm providing an estimation of the input cardinality that has an estimation error of less than 10%.

All of the tested functions performed as expected, thus ensuring a correct implementation.

### 3.2 Evaluating the quality of the hash function

To ensure the quality of the hash function we run an experiment where we determined the distribution of the hash values of  $\rho$  ( $\rho(h(x))$ ) for one million hash values with  $x \in \{1, \dots, 10^6\}$ . The results of this experiment support the assertion that the aforementioned distribution satisfies  $Pr[\rho(y) = i] = 2^{-i}$  for all  $i$  from 1 to  $k$  for random  $y \in \{0, 1\}$ . To check the performance of the function, we stored the results in a list and visualized it on a plot. The following list stores a fraction of  $x$ -es that had a resulting  $\rho$  value - 0.5 at  $\rho = 1$ , 0.25 at  $\rho = 2$  etc.:

[0.5, 0.25, 0.125, 0.0625, 0.03125, 0.01562, 0.00781, 0.00391, 0.00195, 0.00098]. It is easy to notice that the number of  $x$  decreases by half on each next  $\rho$ , what corresponds to the desired probability  $Pr[\rho(y) = i] = 2^{-i}$ .

Figure 1 depicts the described distribution.




Figure 1: Distribution of  $\rho(h(x))$  for one million hash values

## 4 Experiments

The experiments were made on a Surface book 3 with 1.30GHz 8 cores Intel i7 processor, Memory of 32GiB RAMS. Python 3.8.8 was used. We used python’s standard library Random to create an input generator with a seed function to create lists of distinct 32-bit integers. The input generator method was tested for correctness by converting the list into a set to check that there were no duplicates in the list. Moreover, the seed functionality was tested by running the input generation multiple times with a specific seed.

We ran two experiments with five different values for  $m$  (i.e., 256, 512, 1024, 2048, 4096). The first with 100 repetitions and a random input  $n$  distinct 32-bit integers of size 100,000 for each  $m$  test. The second with 1000 repetitions and a smaller random input of size 10,000. This was done by implementing a benchmark function that takes four inputs: the number of repetitions for the

test, the size of  $n$ , the size of the register  $m$ , and  $m\_fit$  value to make the  $f(x)$  (the function that indexes the input to a register) fit the register size. The benchmark function returns a list with the resulting hyperloglog estimations for the specified  $m$  size. For each  $i$  in a repetition in a run of the benchmark function, the  $i$  value was parsed as a seed in the random input generator to get a new  $n$ -sized random list for each repetition. The seed would ensure that for the next experiment with a new  $m$  size, the test input stayed fixed.




Figure 2: Experiment 1 with 100 repetitions,  $n$  of size 10000. Here all five tests are represented ( $m$  values 256, 512, 1024, 2048, 4096). Starting where the back-most histogram represents  $m$  of 256 all the way through to the front-most histogram representing the  $m$  of 4096.




Figure 3: Experiment 2 with 1000 repetitions,  $n$  of size 10000. Here all five tests are represented ( $m$  values 256, 512, 1024, 2048, 4096). Starting where the back-most histogram represents  $m$  of 256 all the way through to the front-most histogram representing the  $m$  of 4096.

Table 1: Experiment 1: 100 repetitions,  $n$  of size 100,000. Fraction of hyperloglog estimates that fall within 1 and 2 standard deviations from  $n$  (the number of distinct elements). Standard deviation is calculated as  $1.04/\sqrt{m}$ .

$m$	$\pm 1\sigma$	$\pm 2\sigma$
256	0.660000	0.980000
512	0.650000	0.940000
1024	0.580000	0.960000
2048	0.740000	0.960000
4096	0.660000	0.960000

Table 2: Experiment 2: 1000 repetitions,  $n$  of size 10,000. Fraction of hyperloglog estimates that fall within 1 and 2 standard deviations from  $n$  (the number of distinct elements). Standard deviation is calculated as  $1,04/\sqrt{m}$ .

$m$	$\pm 1\sigma$	$\pm 2\sigma$
256	0.682000	0.950000
512	0.683000	0.955000
1024	0.678000	0.953000
2048	0.716000	0.974000
4096	0.260000	0.645000

## 5 Results

For the tests with small  $m$  sizes and  $i = 100$  and  $n = 100,000$ , the HyperLogLog estimations are slightly skewed (skewed right for  $m = 256$  and skewed left for  $m = 512$ ). This imprecision could be explained by the large  $n/m$ . We do not see the same when the repetitions are raised from 100 to 1000 in the second test and the  $n$ -size was lowered to 10,000 in experiment 2. In general, the results of experiment 2 come out with a much more even normal distribution bell curve around the expected  $n$ -distinct elements. Except when  $m$  gets very large in relation to  $n$  (i.e.,  $m = 4096$ ). Here the distribution comes out bi-nominal. This is somewhat understandable as we have the small  $n/m$  value of 2,4, which can result in more empty registers. There is a pattern throughout both experiments (seen both on figure 2 and 3, that shows that the variance in the results becomes smaller as the size of the register  $m$  increases. In other words when we increase  $m$  on a fixed size of  $n$ , we see a gain in accuracy as  $m$  increases. Moreover, common for both experiments (with the exception of experiment 2,  $m = 4096$  with  $n/m = 2,4$ ), we have a normal distribution of the fraction of estimates that lie within  $\pm 1\sigma$  (i.e., 68%) and  $\pm 2\sigma$  (i.e., 95%) as seen in table 1 and 2 when  $\sigma$  is derived in relation to the size of  $m$  ( $\sigma = 1,04/\sqrt{m}$ ).

# Appendices

Results from Experiment 1




Figure 4: Experiment 1 with 100 repetitions,  $n$  of size 100000.  $m=256$






Figure 5: Experiment 1 with 100 repetitions,  $n$  of size 100000.  $m=512$




Figure 6: Experiment 1 with 100 repetitions,  $n$  of size 100000.  $m=1024$




Figure 7: Experiment 1 with 100 repetitions,  $n$  of size 100000.  $m=2048$




Figure 8: Experiment 1 with 100 repetitions,  $n$  of size 100000.  $m=4096$

Results from Experiment 2




Figure 9: Experiment 2 with 1000 repetitions,  $n$  of size 10000.  $m=256$




Figure 10: Experiment 2 with 1000 repetitions,  $n$  of size 10000.  $m=512$




Figure 11: Experiment 2 with 1000 repetitions,  $n$  of size 10000.  $m=1024$




Figure 12: Experiment 2 with 1000 repetitions,  $n$  of size 10000.  $m=2048$




Figure 13: Experiment 2 with 1000 repetitions,  $n$  of size 10000.  $m=4096$